

Rethinking the Arithmetic of Multi-Agent Medical Consultation

Anonymous ACL submission

Abstract

Multi-agent consultation has become the default design pattern in medical language-model systems, on the intuition that a panel of specialists diagnoses better than one physician. That intuition inherits an arithmetic premise from human medicine: pooling judgements helps because the judgements pooled are independent. Does it hold for language models? We measure the pairwise error correlation between panel members across 180 multi-agent configurations and 63,499 episodes, spanning three capability tiers, three medical benchmarks, five architectures and panel sizes $N \in \{1, 3, 5, 7, 9\}$. It is $\varphi = 0.73$, which leaves a nine-member panel with 1.31 **effective independent opinions**. Peer discussion raises the correlation to 0.986, so a six-agent consultation is informationally equivalent to 1.01 physicians at the moment it reaches consensus. Two consequences follow. Panel size predicts collaboration gain at cross-validated $R^2 = -0.006$, because tripling a panel adds a tenth of an opinion. And unanimity rises from 66% to 98% while accuracy given unanimity stays at chance, with confidence rising most on wrong answers, because agreement among copies reads as corroboration. The benefit that survives is confined to a window of task difficulty between 25% and 50% single-doctor accuracy, inside which all 20 significant improvements fall and outside which all 10 significant degradations do. Making medical panels agree is easy. Making them independent is the open problem.

1 Introduction

Ask six AI specialists a hard clinical question and let them confer. They will reach unanimity 98% of the time. They will be wrong half the times they agree. And by the moment they agree, they carry 1.01 independent opinions between them. Those three facts are one fact, and it is the subject of this paper.

Multi-agent “consultation” has become the default design pattern in medical language-model systems. MedAgents (Tang et al., 2024), MDAgents (Kim et al., 2024), MAC (Chen et al., 2025) and TeamMedAgents (Mishra et al., 2026) all instantiate the same clinical intuition. A panel of specialists should diagnose better than one generalist, the reasoning goes, because a tumour board diagnoses better than one physician. The pattern is being deployed faster than it is being measured.

The intuition is not merely an analogy; it is an inheritance. Pooling diagnostic judgements improves accuracy in human medicine for a specific reason, namely that physicians trained differently make different mistakes (Kurvers et al., 2016; Kämmer et al., 2017). Majority voting converts that independence into accuracy. Strip the independence away and the arithmetic stops working, however many voters remain. Nobody has checked whether language-model panels satisfy the premise they borrow.

They do not. Measuring the pairwise error correlation between panel members across 180 multi-agent configurations, we find $\varphi = 0.73$: agents prompted into different medical specialties fail on the same items. A nine-member panel therefore carries 1.31 effective independent opinions, and peer discussion drives the correlation to 0.986, at which point a six-agent consultation is informationally equivalent to 1.01 physicians. The panel does not converge on an answer. It collapses onto one.

Three observable consequences follow, and they are the same fact seen three ways. Panel size predicts nothing, because tripling the panel adds a tenth of an opinion. Consensus carries no information, because it is agreement among copies: unanimity reaches 98% while accuracy given unanimity stays at chance, and confidence rises most on the answers that are wrong. And the benefit that survives is confined to a narrow *window* of task dif-

085 difficulty rather than bounded by a ceiling, roughly
 086 25–50% single-doctor accuracy, inside which ev-
 087 ery significant improvement we measure falls and
 088 outside which every significant degradation does.

089 Establishing this requires a grid large enough
 090 to separate architecture from panel size from
 091 task difficulty: 180 multi-agent configurations
 092 and 63,499 episodes across three capability tiers,
 093 three medical benchmarks, five canonical archi-
 094 tectures, panel sizes $N \in \{1, 3, 5, 7, 9\}$, and
 095 a budget-matched single-model control at every
 096 cell. Medicine makes the measurement possible in
 097 two ways that general-domain agentic benchmarks
 098 cannot. Multiple-choice clinical reasoning uses no
 099 tools, so coordination can be measured without
 100 the confound of tool-call overhead. And expert-
 101 level items drive single-agent accuracy down to
 102 15%, far enough into the difficult regime to ex-
 103 pose the lower edge of the window, which is invis-
 104 ible at benchmark difficulties where single agents
 105 already succeed. Scaling principles for general-
 106 domain agent systems (Kim et al., 2026) sup-
 107 ply the architecture taxonomy and the controlled-
 108 comparison design our grid is built on; the error-
 109 correlation and confidence-calibration measure-
 110 ments that carry our argument are new here.

111 2 Related work

112 **Why pooled judgements help, and what**
 113 **they require.** Collective-intelligence results in
 114 medicine rest on error independence: pooling in-
 115 dependent diagnostic judgements improves accu-
 116 racy (Kurvers et al., 2016; Wolf et al., 2015),
 117 plateauing at roughly five raters (Kämmer et al.,
 118 2017), and groups of nine outperform individu-
 119 als on the Human Dx platform (Barnett et al.,
 120 2019), the clinical instance of a general prin-
 121 ciple (Surowiecki, 2004) that multidisciplinary
 122 tumour boards institutionalise (Radcliffe et al.,
 123 2019). The framing question is not whether two
 124 heads beat one but when (Hautz and Kämmer,
 125 2024). The same literature documents what hap-
 126 pens when independence fails: conformity to a
 127 unanimous majority (Asch, 1956) and groupthink
 128 in deliberating groups (Janis, 1972). Correlated er-
 129 rors across models bound ensemble benefit in gen-
 130 eral (Kim et al., 2025), and multi-agent debate can
 131 converge on confident-but-wrong answers (Zhu
 132 et al., 2026). This paper measures that failure
 133 directly in clinical panels, where diagnostic er-
 134 ror is already common enough in ordinary prac-

tice (Graber, 2013; Singh et al., 2014) that manu-
 factured agreement is a safety concern rather than
 an efficiency one.

Medical language models and multi-agent
systems. Medical question answering became
 a capability benchmark once models matched
 licensing-exam thresholds (Kung et al., 2023; Nori
 et al., 2023; Singhal et al., 2023, 2025), and later
 work moved to conducting encounters (Tu et al.,
 2025; McDuff et al., 2025). Prompting strategy
 carries much of that progress (Nori et al., 2024),
 which is why architectural claims need budget-
 matched controls. On the multi-agent side, MedA-
 gents (Tang et al., 2024) introduced role-played
 specialist collaboration, MDAgents (Kim et al.,
 2024) added adaptive routing, and MAC (Chen
 et al., 2025) and TeamMedAgents (Mishra et al.,
 2026) formalised teamwork components, the latter
 sweeping team sizes 2–5. Every published medi-
 cal ablation stops at $N \leq 5$, none reports a dose-
 response curve, and none runs a budget-matched
 single-model control. Skeptical evaluations find
 that orchestration often fails to beat well-prompted
 single models at substantially higher cost (Zhu
 et al., 2025; Liu et al., 2026; Zhao et al., 2025);
 our independence measurement supplies the mech-
 anism those results lack.

Scaling agents, and scaling inference. Debate
 protocols improve factuality (Du et al., 2024;
 Liang et al., 2024) and frameworks such as
 MetaGPT (Hong et al., 2024) and AutoGen (Wu
 et al., 2024) made the pattern easy to de-
 ploy (Zhang et al., 2024). On how far it scales,
 “more agents is all you need” (Li et al., 2024) re-
 ported monotone gains, while Chen et al. (2024)
 showed the relationship is non-monotone and Li
 et al. (2026) confirmed an inverted-U in agent
 count; Yang et al. (2026) attributed the shape to
 opinion diversity, showing two diverse agents can
 match sixteen homogeneous ones. Kim et al.
 (2026) fit a predictive scaling principle across
 architectures, model families and agentic bench-
 marks, and supply the architecture taxonomy and
 controlled-comparison design our grid is built on.
 Test-time compute offers the alternative to adding
 agents (Kaplan et al., 2020; Hoffmann et al., 2022;
 Brown et al., 2024; Snell et al., 2024), and self-
 consistency (Wang et al., 2023) over chain-of-
 thought (Wei et al., 2022) is the budget-matched
 control we run at every cell.

Confidence and calibration. Our safety analysis depends on whether stated confidence is informative. Models are partially aware of what they know (Kadavath et al., 2022), but elicited confidence is weakly calibrated and sensitive to how it is asked (Xiong et al., 2024; Guo et al., 2017), and clinical evaluations report fluent but unreliable responses (Goodman et al., 2023) with demographic biases that propagate (Omiye et al., 2023; Zack et al., 2024). We add a coordination-specific finding: deliberation degrades the discriminative power of confidence rather than improving it.

3 Agent systems and tasks

3.1 System definition

We describe an agent system as a triple (A, C, Ω) : a set of agents A , a communication graph C , and an aggregation rule Ω . Figure 1 and Table 1 give the five configurations we evaluate, each paired with the clinical role it instantiates. The taxonomy is the one used for general-domain agent systems (Kim et al., 2026); the clinical instantiations are ours. Only the coordination logic differs across architectures. The item rendering, the answer schema, the role prompts and the decoding parameters are shared.

Independent. N specialists receive the item and answer without seeing one another. Aggregation is majority vote, ties broken by mean self-reported confidence and then by option order. A confidence-weighted variant is logged throughout.

Centralized (attending physician). An orchestrator receives the N specialist opinions on a star topology and is instructed to *audit before aggregating*. It must name the single opinion most likely to be wrong and justify that, then either commit to an answer or re-task one named specialty. A re-tasked specialist re-answers with the orchestrator’s critique in hand. The star topology is preserved, since specialists never see each other. This design gives the panel a single point at which errors can be caught before they reach the output, and it is the architecture responsible for both the largest gain and the largest loss in the study.

Decentralized (MDT discussion). N specialists give first opinions, then exchange opinions on a complete graph for up to three rounds, each free to revise. Rounds stop early only on true unanimity, defined as every agent emitting the same *valid*

option. A panel containing unparsed answers is not unanimous.

Hybrid (tiered referral). A generalist answers first with a stated confidence. Below a frozen threshold θ the case is referred to the N -specialist panel; if that panel has no majority the case escalates to full MDT discussion. θ is calibrated once on a 100-item development slice disjoint from every evaluation set and then frozen.

Panel composition. A cheap router (gpt-4.1-nano, temperature 0) ranks the nine most relevant specialties for each item from a fixed 17-specialty roster; a panel of size N is the first N of that ranking. Panels are therefore *nested*: the $N=9$ panel contains the $N=5$ panel, which contains the $N=3$ panel.

3.2 Nested panels as common random numbers

Every agent’s first opinion is keyed on (item, specialty _{i} , seed, i) and never on N or on the architecture. Consequently the first opinions of a panel of size N are a superset of those for any $N' < N$, and all four multi-agent architectures reuse the same first opinions. This is a deliberate common-random-numbers coupling. It reduces the variance of within-item contrasts between architectures and panel sizes, and it changes the estimand: we measure *accuracy as the same routed panel is extended*, not the expected accuracy of independently redrawn panels at each size. Paired McNemar tests on item-level outcomes remain valid under this coupling; regression models must cluster at the item level rather than treat configurations as independent.

3.3 What we measure

Accuracy alone cannot distinguish a panel that pools complementary evidence from one that repeats itself, so we instrument each episode with five coordination quantities and one derived quantity that the rest of the paper turns on.

We measure five coordination quantities. Their names are established in the agent-scaling literature (Kim et al., 2026), which reports headline values without always giving closed forms; the definitions below are the ones we compute.

- **Turns T :** reasoning–response exchanges; one model call is one turn.




Figure 1: **The five architectures we evaluate**, with the clinical role each instantiates. A is the agent set, C the communication graph, Ω the aggregation rule.

Table 1: Architectures, their coordination structure, and the clinical role each instantiates. n is the panel size, k reasoning iterations, r orchestration rounds, d debate rounds.

Architecture	(A, C, Ω)	Cost	Clinical instantiation
Single-agent (SAS)	$\{a\}$	$O(k)$	One physician, zero-shot or CoT
Independent	$\{a_1..a_n\}$, $C = \emptyset$, synthesis	$O(nk)$	Specialists answer blind; majority vote
Centralized	$+a_{\text{orch}}$, star, hierarchical	$O(rnk)$	Attending physician audits, then decides
Decentralized	$\{a_1..a_n\}$, complete, consensus	$O(dnk)$	Multidisciplinary team discussion
Hybrid	$+a_{\text{orch}}$, star + peer, both	$O(rnk+pn)$	Tiered referral: triage $\rightarrow$ panel $\rightarrow$ MDT

- **Message density c :** peer-opinion exposures per turn. Independent has $C = \emptyset$ and therefore $c = 0$ by construction; Centralized exposes opinions only to the orchestrator; Decentralized and Hybrid expose every agent to every other each round.
- **Redundancy R :** mean pairwise Jaccard overlap of rationale tokens among first opinions. High R means the panel is restating one another rather than contributing independent information.
- **Overhead $O\%$:** token cost relative to the single-agent baseline on the same items, $(\text{tok}_{\text{MAS}} - \text{tok}_{\text{SAS}})/\text{tok}_{\text{SAS}}$.
- **Coordination efficiency E_c :** accuracy divided by the token-cost ratio to the single-agent baseline.
- **Error absorption:** $(E_{\text{SAS}} - E_{\text{MAS}})/E_{\text{SAS}}$, where E is the error rate.
- **Error amplification A_e :** the extent to which individual-agent errors reach the final output without inter-agent verification. We measure that failure directly, in which the panel *held* the correct option among its first opinions and still emitted a wrong one, normalised by the marginal probability that a single agent is

wrong. $A_e > 1$ means the system destroys information it already possessed.

Effective independent opinions. Majority voting converts independent judgements into accuracy; correlated judgements it merely repeats. We therefore summarise a panel not by its size but by how many independent opinions it carries. Writing φ for the mean pairwise error correlation between panel positions, measured as the correlation of their per-item error indicators, the effective number of independent opinions in a panel of N agents is

$$N_{\text{eff}} = \frac{N}{1 + (N-1)\varphi}. \quad (1)$$

At $\varphi = 0$ the N opinions are fully independent and $N_{\text{eff}} = N$; at $\varphi = 1$ they are one opinion copied N times and $N_{\text{eff}} = 1$. This is the standard design-effect correction for clustered observations, and it makes the arithmetic premise of consultation measurable rather than assumed. Reporting N_{eff} alongside N turns out to explain why panel size predicts so little (§4.5).

3.4 Clinical tasks and benchmarks

Three benchmarks, 250 items each, sampled with a fixed seed: MedXpertQA-Text (Zuo et al., 2025) (10 options; stratified by the 12 body systems $\times$

3 task types), MedQA-USMLE (Jin et al., 2021) (4 options), and the contamination-screened hard split of MedAgentsBench (Tang et al., 2025) (894 items excluding the medqa_5options subset; the first 250 after stratified shuffling). The primary endpoint is exact match against the dataset’s own gold option key. No LLM judge appears anywhere in the evaluation path.

Every item is additionally tagged easy / medium / hard by the pass rate of three mid-tier models at $k=5$ samples. On MedXpertQA the strata are heavily skewed hard (13 / 62 / 125 on the pilot sample), which is the intended regime.

4 Experiments and results

4.1 Setup

LLMs and capability scaling. Three OpenAI models form a capability ladder, measured rather than assumed (Table 11). The ladder is monotone on both benchmarks and spans single-doctor accuracy from 25% to 55% on MedXpertQA, straddling the previously reported 45% threshold from both sides. A task-grounded capability metric explains more variance than a generic intelligence index (Kim et al., 2026), so we take I to be each model’s single-doctor chain-of-thought accuracy on MedQA.

Budget-matched controls. Every configuration is accompanied by three single-model controls. These are zero-shot, chain-of-thought, and **self-consistency at a matched budget**. The self-consistency pool is built from cached $n=8$ sampling calls (the API maximum), so any k is a prefix of the same samples. Nominal cost is charged as $\lceil k/8 \rceil$ prompts plus k completions, which is what an efficient implementation would actually pay. Panels are compared against the self-consistency cell whose mean per-item *dollar* cost is closest, not the one with a matching sample count. Matching on samples silently favours the panel, because a panel pays the prompt once per agent while self-consistency amortises it across a batch.

Cost accounting is *nominal* throughout. Cached calls are charged at list price, so the economics figures report what a configuration would cost to run, not what our cache made it cost.

Scaling law. We fit the functional form used for general-domain agent systems (Kim et al., 2026),

$$P = \beta_0 + \beta_1 I + \beta_2 I^2 + \beta_3 \log(1+T_{\text{tools}}) + \beta_4 \log(1+n_a) + \beta_5 \log(1+O\%) + \beta_6 c + \beta_7 R + \beta_8 E_c + \beta_9 \log(1+A_e) + \beta_{10} P_{\text{SA}} + \sum_j \beta_j (\text{interactions}), \quad (2)$$

with interactions $I \times E_c$, $A_e \times P_{\text{SA}}$, $R \times n_a$, $c \times I$ and $P_{\text{SA}} \times \log(1+n_a)$. Our tasks use no tools, so $T_{\text{tools}} = 0$ identically and every term containing it vanishes. We fit the $T=0$ slice, which is the point of the study. Standard errors are clustered by benchmark, and cross-validated R^2 is reported over five folds.

Statistics. Paired McNemar tests with exact binomial p -values on identical item sets, Holm-corrected within each family of comparisons. Wilcoxon 95% confidence intervals for all accuracies. Power was calibrated on a 200-item pilot. Median discordance between configurations is 8.5%, giving a minimum detectable difference of 5.7 percentage points at $n=200$ and 2.6 points at $n=1000$ with 80% power. Configurations that fail for infrastructure reasons (rate limits, timeouts) are recorded with an error status and excluded from accuracy. Scoring them as wrong answers would bias call-heavy architectures downward.

4.2 Main results

Across 180 multi-agent configurations and 63,499 episodes we find that collaboration in medicine is governed by task difficulty, not by panel size. Figure 3 gives the full dose–response surface; Table 2 the underlying numbers.

Collaboration pays only inside a window of task difficulty. Binning all 180 multi-agent configurations by the single-doctor baseline P_{SA} of their cell reveals a window rather than a ceiling (Fig. 2):

P_{SA}	configurations	mean gain	fraction positive
< 25%	40	−0.88 pp	45%
25–35%	40	+2.29 pp	92%
35–50%	40	+2.67 pp	85%
50–70%	20	−0.88 pp	50%
> 70%	40	+0.55 pp	68%

Configurations inside the 25–50% band gain +2.48 pp on average, while those outside it lose 0.31 pp (Welch $t = 8.42$, $p < 10^{-6}$, $n = 80$ vs 100). Splitting instead at the 45% figure of Kim et al. (2026) reproduces their direction but explains the data far less well (+1.36 pp below versus +0.07 pp above; $t = 3.46$, $p = 0.0007$). The




Figure 2: **Collaboration pays only inside a window of task difficulty.** (a) Collaboration gain against the single-doctor baseline P_{SA} for every multi-agent configuration; The shaded band marks the 25–50% window; the dashed line is the previously reported 45% capability ceiling. (b) The same data binned; annotations give the fraction of configurations in each bin that improve on the single doctor.




Figure 3: **Agent scaling across model capability and system architecture.** Performance against the task-grounded capability index I (mean single-doctor accuracy across the three benchmarks), for each architecture, on each benchmark. Each architecture is shown at its best-performing panel size. Dashed grey: single-agent baseline. Green arrows give the gain of the best architecture over the single agent at the strongest model.

previously reported ceiling is the window’s *upper* edge. Its lower edge becomes visible only because expert-level medical items drive single-agent accuracy down to 15%. A window is not a weaker version of a ceiling. It identifies two distinct failure modes, and it describes the data better by a factor of 2.4 in t .

The window separates significant help from significant harm, with no overlap. Thirty of the 180 configurations differ significantly from the single doctor (McNemar $p < 0.05$), and they partition perfectly:

	sig. better	sig. worse
inside the window ($n = 80$)	20	0
outside the window ($n = 100$)	0	10

Every configuration that significantly beats the

single doctor lies inside the window. Every configuration that significantly loses to it lies outside. The largest gain is the attending-physician (Centralized) panel at $N=3$: +7.60 pp at $P_{SA} = 41.6\%$ (gpt-5-mini, $p = 2 \times 10^{-5}$), with +6.40 pp at $P_{SA} = 32.0\%$ (gpt-5-nano, $p = 0.0037$). The largest loss is also a Centralized panel, at $P_{SA} = 66.0\%$ with a weak model: −13.60 pp ($p = 10^{-5}$), where an orchestrator too weak to audit its own specialists overrides answers the panel had right. The window is not only a description of where gains concentrate. It is a two-sided boundary, and the harm on the far side is as significant as the benefit within.

Architecture matters more than panel size. Within the window, the ordering is stable: Cen-

Table 2: **Main results.** For each capability tier and benchmark: the single-doctor baseline, then each architecture at its best panel size, and the gain of the best architecture over the single doctor. Bold marks the best architecture in each row. Shaded rows are the four cells whose single-doctor baseline falls inside the 25–50% collaboration window; the two largest gains in the study, +6.0 and +7.6 pp, are both shaded. The complete 180-configuration grid is Table 7.

	Single doctor	Best multi-agent variant (accuracy % / n_a)				Gain
Tier		Independ.	Central.	Decentr.	Hybrid	best
<i>MedXpertQA</i> — expert-level reasoning, 10 options						
T1	15.6	16.0/9	19.2/1	16.8/5	18.0/9	+3.6
T2	32.0	33.8/5	37.2/9	38.0/7	35.2/5	+6.0
T3	41.6	47.2/5	49.2/3	48.0/5	45.2/1	+7.6
<i>MedAgentsBench-hard</i> — items where strong models still fail						
T1	20.4	18.0/1	21.6/7	18.4/5	16.0/9	+1.2
T2	32.4	34.8/5	34.4/3	34.4/9	34.0/5	+2.4
T3	44.4	44.8/9	46.0/7	46.8/9	46.0/1	+2.4
<i>MedQA (USMLE)</i> — licensing-exam questions, 4 options						
T1	66.0	67.2/9	64.0/9	67.6/5	65.6/9	+1.6
T2	88.0	89.6/5	89.6/5	89.6/9	88.8/5	+1.6
T3	91.6	92.4/1	93.2/1	92.8/5	93.6/1	+2.0

■ shaded rows have a single-doctor baseline inside the 25–50% collaboration window

tralized and Decentralized lead, Hybrid trails, and Independent voting is indistinguishable from the single doctor. On gpt-5-nano / MedXpertQA, Independent moves from 31.8% at $N=1$ to 32.6% at $N=9$, which is 0.8 pp for nine times the calls, while Decentralized reaches 38.0% at $N=3$ and stays there. Gains saturate at $N=3$ in every cell where they occur, and no architecture in any cell improves beyond $N=5$.

The best architecture changes with capability tier. The best architecture depends on the tier. At T2 the peak is Decentralized (38.0%, +6.0 pp); at T3 it is Centralized (49.2%, +7.6 pp). Family-dependent architectural preferences have been reported for general-domain agents (Kim et al., 2026), and the pattern here is the medical counterpart. The stronger model exploits an orchestrator’s audit, whereas the weaker model benefits more from peer exchange, whose redundancy compensates for individually unreliable opinions.

4.3 Scaling principles

With accuracy as the response, the fit is degenerate. Fitting that form directly gives $R^2 = 0.997$ (5-fold cross-validated $R^2 = 0.996$), which is an artefact of our design rather than a result. The response is accuracy and one predictor is the single-doctor baseline P_{SA} ; the same items and the same model underlie both terms and only the coordination structure varies, so P_{SA} alone nearly determines accuracy ($\beta_{P_{SA}} = 0.875$, $p < 10^{-4}$). We

therefore refit with the collaboration *gain* as the response, which is what the scaling principle is actually about.

Coordination metrics explain the gain that accuracy alone cannot. With $\text{gain} = P - P_{SA}$ as the response, $R^2 = 0.638$ and 5-fold cross-validated $R^2 = 0.415$. Coordination metrics carry that fit with no tool-count term at all. Coefficients ($n = 225$ configurations including the self-consistency control, benchmark-clustered SEs):

term	$\hat{\beta}$	p
$\log(1 + A_e)$	−0.150	< 0.001
R (redundancy)	+0.082	< 0.001
c (message density)	+0.006	< 0.001
$\log(1 + O\%)$	+0.005	0.024
$P_{SA} \times \log(1 + n_a)$	+0.027	0.048
$\log(1 + n_a)$	−0.003	0.701
P_{SA}	−0.113	0.051
P_{SA}^2	+0.037	0.390

Error amplification is the dominant penalty; redundancy and message density carry small positive effects; overhead is not penalised once $T=0$; and neither panel size nor P_{SA} has a main effect.

$T_{\text{tools}} = 0$ throughout, so the $\log(1 + T)$ main effect and the $O\% \times T$ and $E_c \times T$ interactions vanish identically. Coordination metrics remain highly significant without them, which establishes that the coordination effect is not a re-description of tool overhead. The coefficient on $\log(1 + O\%)$ turns *positive* here. With no tools to coordinate over, spending more tokens on deliberation is mildly beneficial rather than costly.

The decision boundary is not identifiable in the tool-free slice. The published threshold is derived as a coefficient ratio, $P_{SA}^* = \hat{\beta}_4/\hat{\beta}_{17} \approx 0.063/0.408 = 0.154$ in standardised units, or ≈ 0.45 after denormalisation (Kim et al., 2026). In our fit $\hat{\beta}_{\log(1+n_a)} = +0.0005$ ($p = 0.95$) and $\hat{\beta}_{P_{SA} \times \log(1+n_a)} = +0.0074$ ($p = 0.50$). The ratio, 0.064, rests on two null coefficients and carries no information. This is a statement about functional form, not a failed replication. A linear interaction between baseline and log panel size can express a monotone ceiling but cannot express a window. The binned analysis (§4.2) shows the relationship is non-monotone in P_{SA} , and a window is what the medical data require.

Robustness of the fit. Our fit carries I and I^2 with large coefficients of opposite sign (-0.231 and $+0.228$), but bootstrap resampling ($n = 1000$) gives them standard errors of 0.94 and 0.54. With three capability tiers the quadratic is unidentifiable and we do not interpret it. Every other large coefficient is stable: $\log(1 + A_e)$ has bootstrap SE 0.024 and P_{SA} has 0.044.

Residuals pass a normality test (Shapiro–Wilk $p = 0.207$) but fail homoscedasticity decisively (Breusch–Pagan $p = 0.0001$). All standard errors in this paper are therefore heteroscedasticity-robust (HC3), clustered by benchmark where more than two clusters exist. Residual standard error is $\hat{\sigma} = 0.017$, reflecting that collaboration gains in medicine are small in absolute terms. Regularised alternatives do not improve prediction: Lasso (10-fold CV for λ) retains 9 of 11 predictors at CV $R^2 = 0.382$ and Ridge reaches 0.144, both below the unregularised 0.415, so we retain the full specification.

Nested model comparison isolates where the explanatory power lives (5-fold CV R^2): the full model reaches 0.415; capability alone 0.196; coordination metrics alone 0.093; the single-doctor baseline alone 0.043; and **panel size alone** -0.006 , worse than predicting the mean. Any account of medical consultation that treats “how many agents” as the design variable is fitting noise.

The scaling principle selects an architecture, given a case mix. The fitted model converts into a design rule. For a clinical deployment characterised by its single-doctor baseline P_{SA} , expected collaboration gain is positive only for $P_{SA} \in [25\%, 50\%]$, is maximised by centralized coordi-

nation at $n_a = 3$, and is negative below 25%. Three archetypes make this concrete. *Board-style examination triage* ($P_{SA} \approx 0.90$, MedQA-like): use a single model; no architecture we tested returns a significant gain and every panel costs 2–5 $\times$ more. *Specialist referral review* ($P_{SA} \approx 0.35$, MedXpertQA-like): use an attending-physician panel at $n_a = 3$; expected gain +6 to +8 pp at a 6.7 $\times$ cost premium. *Frontier diagnostic difficulty* ($P_{SA} < 0.25$, MedAgentsBench-like at weak tiers): use a single model and spend the budget on a stronger one; panels lose 2.8 pp here because a panel of unreliable specialists reinforces its own errors.

Error amplification, not panel size, dominates the fit. $\log(1 + A_e)$ carries the largest coefficient of any coordination metric (-0.150 , $p < 10^{-4}$). A_e measures how often a panel held the correct option among its first opinions and still emitted a wrong one. Configurations that discard information they already possess are exactly the configurations that fail to beat the single doctor. The same mechanism has been proposed from the opposite direction for general-domain agents (Kim et al., 2026); our measurement supplies the configuration-level evidence.

$R \times n_a$ is positive ($\hat{\beta} = +0.006$, $p < 0.001$). In clinical panels, specialists restating one another is a weak form of verification rather than wasted capacity.

4.4 Coordination efficiency, error dynamics, and information transfer

Coordination is cheap in turns and uninformative in messages. Fitting $T = a(n + 0.5)^b$ gives $a = 0.85$, $b = 0.770$, $R^2 = 0.299$. The exponent is sub-linear, so turns grow more slowly than agents. A clinical panel answering a closed question terminates on unanimity rather than negotiating open-ended tool use. Turn counts by architecture are Decentralized 7.5, Centralized 6.5, Independent 5.0, Hybrid 1.9, against a single-agent 1.4, a range of 1.4 to 5.3 $\times$. The hard resource ceiling they infer from super-linear growth does not arise in tool-free consultation.

$S = 0.462 - 0.045 \ln(c)$, $R^2 = 0.056$: message density explains almost none of the variance in accuracy. Our Centralized panels run at $c = 0.99$ and Decentralized at $c = 0.79$, well past the $c^* = 0.39$ plateau reported elsewhere, with no accuracy penalty. Once a panel is communicating at




Figure 4: **Coordination dynamics.** (a) Turn count against agent count, with our fit and the published general-domain power law for comparison; medical panels converse an order of magnitude less than that law predicts. (b) Performance against message density, with a reference plateau marked at $c^* = 0.39$; performance is uncorrelated with how much the panel says. (c) Error amplification A_e , the share of information a round discards, by architecture. (d) Error absorption, the share of individually wrong opinions the panel repairs, against the published 22.7% reference.

all, communicating more neither helps nor hurts.

Absorption $(E_{\text{SAS}} - E_{\text{MAS}})/E_{\text{SAS}}$ is positive for every architecture: Decentralized +4.2%, Hybrid +2.7%, Independent +1.9%, Centralized +1.6%. Absorption is small in absolute terms and the ordering is not the expected one: peer discussion absorbs most, orchestration least.

Coordination failure is structurally impossible without a channel. We classify first-round errors into four categories, following the multi-agent failure taxonomy of Cemri et al. (2025):

	Indep.	Centr.	Decentr.	Hybrid
Logical contradiction	6.3%	6.3%	6.3%	3.4%
Numerical drift	4.5%	4.7%	4.8%	3.1%
Context omission	13.0%	13.3%	13.4%	6.8%
Coordination failure	0.0%	3.0%	1.2%	0.0%

Coordination failure is **0.0%** under Independent coordination, and necessarily so: with $C = \emptyset$ there is no channel on which it could occur. Hybrid is also at 0.0% because its triage gate de-

clines to convene a panel on most items. The magnitudes do not replicate: peer verification is reported to cut logical contradiction from 16.8% to 11.5%, whereas our first-round contradiction rate is flat at 6.3% across communicating and non-communicating architectures. Numerical drift is an order of magnitude rarer than in tool-using tasks (3–5% against 19–26%), as expected when no arithmetic is cascaded.

Information gain is positive but does not buy accuracy. Information gain, the reduction in posterior variance of the success indicator between the first and final round (§D), is +0.90 for Decentralized, +1.46 for Hybrid and +0.24 for Centralized. Panels converge, and they converge most where the architecture communicates most. Convergence occurs whether or not the consensus is correct: $\Delta\mathcal{I}$ does not predict collaboration gain ($r = 0.055$, $p = 0.59$). Medical consultation is information-limited rather than information-transferring.

Architecture rankings do not generalise across benchmarks. Kendall’s τ between the architecture rankings of any two benchmarks averages $+0.067$ (pairwise -0.200 , $+0.400$, 0.000 ; none significant). Leave-one-benchmark-out cross-validation gives $R^2 = -0.015$: a model fitted on two medical benchmarks predicts the third no better than the mean. Coordination principles that transcend task structure in general-domain agentic work do not transcend benchmark in medicine. This is the practical counterpart of the window: architecture must be selected per difficulty regime, and difficulty regime does not coincide with domain.

The architecture that wins on accuracy is the one that costs most. Successes per 1,000 tokens: Hybrid 0.80, self-consistency 0.47, Independent 0.39, Decentralized 0.33, Centralized 0.19. Centralized is $4.3\times$ less token-efficient than the cheapest configuration, and the architecture responsible for every significant gain is the one that costs most per correct answer. Hybrid is the most efficient, because its triage gate declines to escalate most cases and therefore rarely pays for a panel at all.

On MedXpertQA at gpt-5-nano, Decentralized at $N=3$ reaches 38.0% while self-consistency at the closest matched dollar cost reaches 36.0%. Outside the window the control wins or ties: on MedQA at the same tier, self-consistency at $k=15$ (88.4%) is within noise of the best panel (89.6%) at lower cost, and on MedAgentsBench at T1 self-consistency (20.0%) beats every architecture. Coordination buys something beyond repetition, and only where the window says it should.

4.5 The arithmetic premise of consultation

Majority voting has value only when the votes are independent. Every result above is a consequence of how badly that premise fails.

A nine-member panel carries 1.3 independent opinions. Across all architectures at $n_a \geq 3$, the mean pairwise error correlation is $\varphi = 0.73$. By Eq. 1 this leaves $N_{\text{eff}} = 1.21$ at $N=3$, 1.27 at $N=5$ and 1.31 at $N=9$. Tripling the panel adds 0.10 of an independent opinion. Specialty role prompts do not make one model into several physicians; they make it into one physician wearing several labels.

Discussion drives the effective panel size to one. Peer discussion raises φ from 0.751 before the ex-

Table 3: **Effective independent opinions.** φ is the mean pairwise error correlation between panel positions; $N_{\text{eff}} = N/(1 + (N - 1)\varphi)$. Neither growing the panel nor changing how it is wired moves φ . The shaded row is the single exception: Hybrid’s triage gate declines to convene a panel on most items, so its measured pairs come from the harder subset it does escalate.

Architecture	N	φ	N_{eff}	info. used
<i>Growing one panel: tripling N buys 0.10 of an opinion</i>				
Independent	3	0.743	1.21	40%
Independent	5	0.731	1.27	25%
Independent	9	0.734	1.31	15%
<i>Changing the architecture at $N=9$: no effect either</i>				
Centralized	9	0.735	1.31	15%
Decentralized	9	0.735	1.31	15%
Hybrid	9	0.493	1.82	20%

Table 4: **Clinical safety.** Unanimity is measured on the last round with two or more agents; $P(\text{wr.} | \text{un.})$ is the chance a fully agreeing panel is wrong; the last column is the confidence gap between correct and incorrect answers, the signal a clinician would use to discount an answer. The first row is the single doctor, for reference. Decentralized coordination buys the strongest appearance of agreement and the weakest usable signal at the same time.

Architecture	unan.	$P(\text{wr.} \text{un.})$	conf. gap
Single doctor	—	—	3.3
Independent	66.2%	46.5%	2.8
Centralized	68.6%	44.6%	2.7
Decentralized	98.3%	49.9%	1.8
Hybrid	43.9%	59.7%	2.1

the shaded row agrees almost always and is wrong half the time it agrees

change to **0.986** after it (Fig. 5) (paired $t = 20.9$, $p = 8 \times 10^{-7}$). At $\varphi = 0.986$ a six-agent panel has $N_{\text{eff}} = 1.01$. A multidisciplinary consultation, at the moment it reaches consensus, is informationally equivalent to one physician, and the vote taken at that moment is a vote among copies. This is the mechanism behind every safety result in §4.6: unanimity rises to 98% because the opinions have merged, not because they have converged on truth; the error rate given unanimity stays at a coin flip because merging does not add information; and confidence rises on wrong answers because N agents agreeing looks like corroboration when it is repetition.

The nested model comparison found panel size alone at cross-validated $R^2 = -0.006$. The independence measurement explains why that is not a null result but a necessary one: over $N \in$




Figure 5: **Adding agents adds almost no information, and discussion removes what is left.** (a) Effective independent opinions $N_{\text{eff}} = N/(1 + (N - 1)\varphi)$ against panel size, on a logarithmic scale. Dotted curves are contours of fixed error correlation φ ; the solid line at $N_{\text{eff}}=1$ is a panel worth one physician. Measured panels track the $\varphi \approx 0.75$ contour, so nine members carry barely more than one opinion. (b) φ for every one of the 120 grid configurations, one point per configuration, with the median marked. The value is a near-constant of the setting rather than an artefact of any one cell, Hybrid routing is the only architecture that lowers it, and peer discussion drives it to 0.99.

$\{3, \dots, 9\}$, N_{eff} moves from 1.21 to 1.31. The regressor that matters barely varies, so the regressor we can control cannot predict. Any design effort spent on how many agents to deploy is spent on the wrong variable.

What independence does not explain. Error correlation accounts for the size and consensus results; it does not account for the window. φ correlates only weakly with the single-doctor baseline ($r = -0.215$, $p = 0.018$: harder items do carry more correlated errors) and not at all with collaboration gain ($r = -0.083$, $p = 0.37$), and a model containing φ alone reaches $R^2 = 0.007$. The window and the independence collapse are two separate facts about medical consultation: the first says when a panel can help, the second says why it cannot help much, and why its agreement should not be read as evidence.

4.6 Clinical safety: what consultation does to confidence

Accuracy is not the only thing coordination changes. In medicine the more consequential question is what it does to the *signals a clinician would act on*. Three measurements, all pointing the same way.

Self-reported confidence barely discriminates right from wrong, and discussion makes it worse. Across all configurations, mean confidence on correct answers exceeds that on wrong

answers by only 3.3 points for a single doctor (86.0 vs 82.6; AUC = 0.597; Fig. 6). Peer discussion compresses that gap by 45%, to 1.8 points (88.4 vs 86.6; AUC = 0.577). The compression is not a rescaling: discussion raises confidence on *wrong* answers by 4.0 points while raising it on correct answers by 2.4. A panel that has deliberated is more sure of itself, and disproportionately so when it is mistaken. This is consistent with the threshold calibration we ran before the main experiments, where the best achievable referral gate on stated confidence reached a Youden J of only 0.071.

Discussion manufactures consensus without manufacturing correctness. Unanimity, defined as every panelist emitting the same valid option in the last round with two or more agents, rises from 66.2% under independent voting to 98.3% under peer discussion. Conditional accuracy does not follow: $P(\text{wrong} \mid \text{unanimous})$ is 46.5% for Independent and 49.9% for Decentralized. A unanimous medical panel is wrong about half the time in both cases; discussion simply produces far more unanimous panels. Centralized coordination is the only architecture that lowers the conditional error rate (44.6%), which is the error-containment role the orchestrator is supposed to play.

The most trustworthy-looking output is the least trustworthy. Taken together: consultation




Figure 6: **Consultation degrades the signal a clinician would use to catch its errors.** (a) Overconfidence, defined as stated confidence minus the accuracy actually achieved at that confidence, against stated confidence. Perfect calibration is the solid line at zero; every architecture sits 25–43 pp above it at every confidence level, so panels are not merely uninformative about their own reliability but systematically overconfident. (b) The confidence gap between correct and incorrect answers, which is the quantity a reader would use to discount an answer. Peer discussion nearly halves it.

drives panels to near-total agreement, leaves the error rate given agreement at a coin flip, and raises confidence most on the answers that are wrong while shrinking the confidence gap that would let a reader tell the difference. The configuration a clinician would most readily act on, a unanimous and high-confidence multidisciplinary recommendation, carries essentially the same error rate as a divided one, and advertises itself more strongly. Accuracy gains of 6–8 pp inside the window are real, but they are purchased alongside a degradation in the calibration of the signal that would let a clinician catch the remaining errors.

5 Conclusion

Consultation works in human medicine because physicians fail differently. Our measurements say that language-model panels do not. Agents prompted into different specialties make the same mistakes ($\varphi = 0.73$), and discussion pushes them into near-perfect agreement ($\varphi = 0.986$), leaving a six-member consultation with 1.01 effective independent opinions. Every other result in this paper is a shadow of that one.

Panel size cannot matter, because it does not change what the panel knows. Across $N \in \{3, \dots, 9\}$ the effective number of opinions moves from 1.21 to 1.31, and panel size alone predicts collaboration gain at cross-validated $R^2 = -0.006$. Consensus cannot be trusted, because it is agreement among copies. Unanimity reaches 98%

while the error rate given unanimity stays at 50%, and the confidence signal that would let a clinician catch the difference is compressed by 45%. What remains is a narrow window of task difficulty, roughly 25–50% single-doctor accuracy. Inside it, coordination recovers enough complementary information to pay for itself. Outside it, coordination reliably does not.

Two design conclusions follow for medical deployments. First, the variable worth controlling is not how many agents but whether a competent one audits the others before the system commits: centralized coordination with a capable orchestrator produces the largest gain we measure, and the same architecture with a weak orchestrator produces the largest loss. Specialist capability matters $4.1\times$ more than the orchestrator’s, so capability belongs on the agents doing the reasoning. Second, the deployment decision should be made by measuring the single-model baseline on the intended case mix before choosing an architecture at all. Outside the window, a budget-matched self-consistency control matches or beats every architecture we tested, at a fraction of the cost.

The broader caution is about what a consulting AI system advertises. A unanimous, high-confidence multidisciplinary recommendation is the output a clinician is most likely to act on, and in these experiments it is the output whose reliability is least legible from its own signals. Making panels agree is easy. Making them independent is

the open problem.

6 Limitations and future works

Our tasks are multiple-choice examination items, not clinical encounters. They carry no history taking, no test ordering, no longitudinal follow-up, and the exact-match endpoint rewards the option-selection component of diagnosis only. The $T = 0$ property that makes the design clean also bounds what it licenses. Results here do not extrapolate to tool-using clinical agents.

All models come from one vendor, so the heterogeneity arm mixes separately trained OpenAI families rather than distinct vendors. This bears directly on our central claim: cross-vendor panels may decorrelate more than $\varphi = 0.73$, and the effective panel size we report is therefore a lower bound on what heterogeneous deployments could achieve.

Panels are nested by construction (§2.2), so the reported curves estimate accuracy as a routed panel is extended, not the expected accuracy of independently redrawn panels at each size. The coupling reduces variance on the architecture contrasts we care about and does not bias within-item paired tests, but a reader interested in the independent-panel estimand should read the curves accordingly.

The specialty router ranks by relevance, so larger panels append progressively less relevant specialists. “More experts” and “weaker experts” are therefore entangled in the specialty-role condition. We separate them with the generic-internist control, in which every panelist is an undifferentiated internist and panel size is the only thing that varies.

Difficulty strata are defined by the pass rate of mid-tier models, some of which also appear in the evaluation ladder. This inflates the level of the hard-stratum results but not the within-item slope across panel sizes, which is computed on identical items.

References

Solomon E. Asch. 1956. Studies of independence and conformity: A minority of one against a unanimous majority. *Psychological Monographs*, 70(9).

Michael L. Barnett, Dhruv Boddupalli, Shantanu Nundy, and David W. Bates. 2019. Comparative accuracy of diagnosis by collective intelligence of

multiple physicians vs individual physicians. *JAMA Network Open*, 2(3).

Bradley Brown, Jordan Juravsky, Ryan Ehrlich, et al. 2024. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv:2407.21787*.

Mert Cemri et al. 2025. Why do multi-agent llm systems fail? *arXiv*.

Lingjiao Chen, Jared Quincy Davis, Boris Hanin, Peter Bailis, Ion Stoica, Matei Zaharia, and James Zou. 2024. Are more llm calls all you need? towards the scaling properties of compound ai systems. In *NeurIPS*.

Xiangru Chen et al. 2025. A multi-agent conversational framework for medical decision making. *npj Digital Medicine*.

Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. 2024. Improving factuality and reasoning in language models through multiagent debate. In *ICML*.

Rachel S. Goodman, J. Randall Patrinely, Cosby A. Stone, et al. 2023. Accuracy and reliability of chatbot responses to physician questions. *JAMA Network Open*, 6(10).

Mark L. Graber. 2013. The incidence of diagnostic error in medicine. *BMJ Quality & Safety*, 22:ii21–ii27.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On calibration of modern neural networks and llm confidence. *ICML*.

Wolf E. Hautz and Juliane E. Kämmer. 2024. Whether two heads are better than one is the wrong question. *Diagnosis*.

Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, et al. 2022. Training compute-optimal large language models. *NeurIPS*.

Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, et al. 2024. Metagtpt: Meta programming for a multi-agent collaborative framework. In *ICLR*.

Irving L. Janis. 1972. Victims of groupthink: A psychological study of foreign-policy decisions and fiascoes. *Houghton Mifflin*.

Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. 2021. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *Applied Sciences*, 11(14).

Saurav Kadavath, Tom Conerly, Amanda Askell, et al. 2022. Language models (mostly) know what they know. In *arXiv:2207.05221*.

926	Juliane E. Kämmer, Wolf E. Hautz, Stefan M. Herzog, Olga Kunina-Habenicht, and Ralf H. J. M. Kurvers. 2017. The potential of collective intelligence in emergency medicine: Pooling medical students' independent decisions improves diagnostic performance. <i>Medical Decision Making</i> , 37(6):715–724.	980
927		981
928		982
929		983
930		
931		984
932		985
933	Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, et al. 2020. Scaling laws for neural language models. <i>arXiv:2001.08361</i> .	986
934		987
935		
936	Elliot Kim, Avi Garg, Kenny Peng, and Nikhil Garg. 2025. Correlated errors in large language models. In <i>ICML</i> .	988
937		989
938		990
939	Yubin Kim, Ken Gu, Chanwoo Park, Chunjong Park, Samuel Schmidgall, A. Ali Heydari, Yao Yan, Zhihan Zhang, Yuchen Zhuang, Yun Liu, Mark Malhotra, Paul Pu Liang, Hae Won Park, Yuzhe Yang, Xuhai Xu, Yilun Du, Shwetak Patel, Tim Althoff, Daniel McDuff, and Xin Liu. 2026. Capable language models can outgrow the benefits of collaboration . <i>Nature Machine Intelligence</i> , 8:1157–1172. Preprint: <i>arXiv:2512.08296</i> , <i>Towards a Science of Scaling Agent Systems</i> .	991
940		992
941		993
942		994
943		995
944		
945		996
946		997
947		998
948		
949	Yubin Kim, Chanwoo Park, Hyewon Jeong, Yik Siu Chan, Xuhai Xu, Daniel McDuff, Hyeonhoon Lee, Marzyeh Ghassemi, Cynthia Breazeal, and Hae Won Park. 2024. Mdagents: An adaptive collaboration of llms for medical decision-making. In <i>NeurIPS</i> .	999
950		1000
951		1001
952		
953		1002
954	Tiffany H. Kung, Morgan Cheatham, Arielle Medenilla, et al. 2023. Performance of chatgpt on usmle. <i>PLOS Digital Health</i> , 2(2).	1003
955		1004
956		1005
957	Ralf H. J. M. Kurvers, Stefan M. Herzog, Ralph Herzig, Jens Krause, Patricia A. Carney, Andy Bogart, Giuseppe Argenziano, Iris Zalaudek, and Max Wolf. 2016. Boosting medical diagnostics by pooling independent judgments. <i>PNAS</i> , 113(31):8777–8782.	1006
958		1007
959		1008
960		1009
961		
962	Jialing Li, Zhouhong Gu, Yin Cai, and Hongwei Feng. 2026. Scaling behavior of single llm-driven multi-agent systems. <i>arXiv:2606.00655</i> .	1010
963		1011
964		1012
965	Junyou Li, Qin Zhang, Yangbin Yu, Qiang Fu, and Deheng Ye. 2024. More agents is all you need. <i>TMLR</i> .	1013
966		
967	Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. 2024. Encouraging divergent thinking in large language models through multi-agent debate. In <i>EMNLP</i> .	1014
968		1015
969		
970		1016
971		1017
972	Y. Liu, Zunamys I. Carrero, X. Jiang, et al. 2026. Benchmarking large language model-based agent systems for clinical decision tasks . <i>npj Digital Medicine</i> .	1018
973		1019
974		1020
975		
976	Daniel McDuff, Mike Schaekermann, Tao Tu, Anil Palepu, et al. 2025. Towards accurate differential diagnosis with large language models. <i>Nature</i> , 642:451–457.	1021
977		1022
978		1023
979		1024
	Pranav Pushkar Mishra, Mohammad Arvan, and Mohan Zalake. 2026. Teammedagents: Pareto-efficient multi-agent medical reasoning through teamwork theory. <i>arXiv:2508.08115</i> .	1025
		1026
		1027
	Harsha Nori, Nicholas King, Scott Mayer McKinney, Dean Carignan, and Eric Horvitz. 2023. Capabilities of gpt-4 on medical challenge problems. <i>arXiv:2303.13375</i> .	
		1028
	Harsha Nori, Naoto Usuyama, Nicholas King, et al. 2024. From medprompt to ol: Exploration of run-time strategies for medical challenge problems. <i>arXiv:2411.03590</i> .	1029
		1030
		1031
		1032
	Jesutofunmi A. Omiye, Jenna C. Lester, Simon Spichak, Veronica Rotemberg, and Roxana Daneshjou. 2023. Large language models propagate race-based medicine. <i>npj Digital Medicine</i> , 6.	
		1033
	Emma Radcliffe et al. 2019. Multidisciplinary tumour boards and their impact on cancer care. <i>The Lancet Oncology</i> .	1034
		1035
	Hardeep Singh, Ashley N. D. Meyer, and Eric J. Thomas. 2014. The frequency of diagnostic errors in outpatient care. <i>BMJ Quality & Safety</i> , 23(9).	
		1036
	Karan Singhal, Shekoofeh Azizi, Tao Tu, S. Sara Mahdavi, Jason Wei, Hyung Won Chung, et al. 2023. Large language models encode clinical knowledge. <i>Nature</i> , 620:172–180.	1037
		1038
	Karan Singhal, Tao Tu, Juraj Gottweis, Rory Sayres, et al. 2025. Toward expert-level medical question answering with large language models. <i>Nature Medicine</i> , 31:943–950.	1039
		1040
	Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. Scaling llm test-time compute optimally can be more effective than scaling model parameters. <i>arXiv:2408.03314</i> .	1041
		1042
	James Surowiecki. 2004. The wisdom of crowds. <i>Doubleday</i> .	1043
		1044
	Xiangru Tang, Anni Zou, Zhuosheng Zhang, Ziming Li, Yilun Zhao, Xingyao Zhang, Arman Cohan, and Mark Gerstein. 2024. Medagents: Large language models as collaborators for zero-shot medical reasoning. In <i>Findings of ACL</i> .	1045
		1046
	Xiangru Tang et al. 2025. Medagentsbench: Benchmarking thinking models and agent frameworks for complex medical reasoning. <i>Patterns / arXiv:2503.07459</i> .	1047
		1048
	Tao Tu, Anil Palepu, Mike Schaekermann, Khaled Saab, et al. 2025. Towards conversational diagnostic artificial intelligence. <i>Nature</i> , 642:442–450.	1049
		1050
	Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models. In <i>ICLR</i> .	1051
		1052

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *NeurIPS*.

Max Wolf, Jens Krause, Patricia A. Carney, Andy Bogart, and Ralf H. J. M. Kurvers. 2015. Collective intelligence meets medical decision-making. *PLOS ONE*, 10(8).

Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, et al. 2024. Autogen: Enabling next-gen llm applications via multi-agent conversation. In *COLM*.

Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2024. Can llms express their uncertainty? an empirical evaluation of confidence elicitation in llms. In *ICLR*.

Yingxuan Yang, Chengrui Qu, Muning Wen, Laixi Shi, Ying Wen, Weinan Zhang, Adam Wierman, and Shangding Gu. 2026. Understanding agent scaling in llm-based multi-agent systems via diversity. *arXiv:2602.03794*.

Travis Zack, Eric Lehman, Mirac Suzgun, et al. 2024. Assessing the potential of gpt-4 to perpetuate racial and gender biases in health care. *The Lancet Digital Health*, 6(1).

Jintian Zhang, Xin Xu, Ningyu Zhang, Shumin Deng, Chuanqi Tan, and Huajun Chen. 2024. Exploring collaboration mechanisms for llm agents. *ACL*.

Huiya Zhao, Yinghao Zhu, Zixiang Wang, Yasha Wang, Junyi Gao, and Liantao Ma. 2025. Confagents: A conformal-guided multi-agent framework for cost-efficient medical diagnosis. *arXiv:2508.04915*.

Xiaochen Zhu, Caiqi Zhang, Yizhou Chi, Tom Stafford, Nigel Collier, and Andreas Vlachos. 2026. Demystifying multi-agent debate: The role of confidence and diversity. *arXiv:2601.19921*.

Yinghao Zhu, Ziyi He, Haoran Hu, Xiaochen Zheng, Xichen Zhang, Zixiang Wang, Junyi Gao, Liantao Ma, and Lequan Yu. 2025. Medagentboard: Benchmarking multi-agent collaboration with conventional methods for diverse medical tasks. *arXiv:2505.12371*.

Yuxin Zuo, Shang Qu, Yifei Li, Zhangren Chen, Xuekai Zhu, Ermo Hua, Kaiyan Zhang, Ning Ding, and Bowen Zhou. 2025. Medxpertqa: Benchmarking expert-level medical reasoning and understanding. In *ICML*.

A Model capability index

Prior work scores models on a published Intelligence Index and reports that a *task-grounded* capability metric explains more variance ($R^2 = 0.413$) than the generic one ($R^2 = 0.373$). We therefore use the task-grounded form throughout: I

is a model’s mean single-doctor chain-of-thought accuracy across the three benchmarks. Table 5 gives the components. The three tiers span $I = 34.0$ to 59.2 , and their single-benchmark baselines span 15.6% to 91.6% , crossing the 45% region from both sides. That crossing is what makes the capability-ceiling test possible at all.

B Domain complexity

B.1 Complexity metric construction

Domain complexity $D \in [0, 1]$ is the arithmetic mean of three complementary measures (Kim et al., 2026): the **performance ceiling** $1 - p_{\max}$, where $p_{\max}$ is the highest performance any evaluated system reaches; the **coefficient of variation** σ/μ of performance across all configurations, a scale-invariant measure of relative spread; and the **best-model baseline** $1 - p_{\text{best}}$, where p_{best} is the strongest single-model performance on the dataset.

B.2 Domain characterisation

Table 6 gives the components. MedQA is a low-complexity domain ($D = 0.097$); MedXpertQA and MedAgentsBench-hard are high-complexity and nearly indistinguishable ($D = 0.489$ and 0.478).

B.3 The critical threshold does not transfer

A critical complexity threshold at $D \approx 0.40$ has been reported for general-domain agents (Kim et al., 2026): below it multi-agent architectures yield net positive returns, above it coordination overhead consumes resources otherwise spent on reasoning. Our data do not support a benchmark-level threshold of any value. MedXpertQA ($D = 0.489$) and MedAgentsBench-hard ($D = 0.478$) differ by 0.011 in complexity yet by 3.5 pp in mean multi-agent gain ($+3.09$ pp versus -0.37 pp), giving opposite signs at effectively the same D . A benchmark-level statistic cannot separate them because the quantity that does separate them varies *within* each benchmark: on MedAgentsBench-hard the per-tier gains run -2.78 , $+1.02$, $+0.66$ pp as P_{SA} moves from 20.4% through 32.4% to 44.4% , entering the window between the first tier and the second. Domain complexity is the right idea at the wrong granularity for medicine. The configuration-level baseline P_{SA} is what carries the signal.

Table 5: Capability index and its components (single-doctor CoT accuracy, %).

Model	MedXpertQA	MedAgentsBench	MedQA	I (mean)	Reasoning
gpt-4.1-nano	15.6	20.4	66.0	34.0	—
gpt-5-nano	32.0	32.4	88.0	50.8	effort=low
gpt-5-mini	41.6	44.4	91.6	59.2	effort=low

shaded tiers straddle the 45% ceiling reported for general-domain agents

Table 6: Domain complexity of the three medical benchmarks.

Benchmark	$1 - p_{\max}$	σ/μ	$1 - p_{\text{best}}$	D	mean MAS gain
MedXpertQA	0.508	0.376	0.584	0.489	+3.09 pp
MedAgentsBench-hard	0.532	0.346	0.556	0.478	−0.37 pp
MedQA (USMLE)	0.060	0.147	0.084	0.097	+0.07 pp

shaded benchmarks are the two whose difficulty places them inside the window

C Datasets

MedXpertQA-Text (Zuo et al., 2025): 2,450 expert-level items, ten options each, tagged by 12 body systems and 3 task types. We draw a stratified sample interleaved across (body system $\times$ task) so that any prefix is approximately balanced, then take the first 250. **MedQA-USMLE** (Jin et al., 2021): 1,273 four-option test items; randomly shuffled with a fixed seed, first 250. **MedAgentsBench-hard** (Tang et al., 2025): the contamination-screened hard split, 894 items after excluding the medqa_5options subset, stratified-interleaved across the ten source subsets, first 250. All samples are nested prefixes, so a larger sample re-uses every call of a smaller one.

Item difficulty is additionally tagged easy/medium/hard by the pass rate of three mid-tier models at $k=5$; on the MedXpertQA sample this yields 13/62/125 items.

D Implementation details

D.1 Technical infrastructure

All models are accessed through one OpenAI-compatible client with a content-addressed disk cache, exponential backoff with rate-limit-aware retry, a per-call JSONL ledger recording tokens, latency and cost, and a cross-process spend journal guarded by an OS file lock. Every LLM call is cached on the full request (model, system, messages, temperature, seed, max tokens, effort, JSON mode, tag), so a re-run is a replay.

D.2 Agent configuration

Single agents answer in one call. Independent panels deploy N specialists with synthesis-only aggregation. Centralized systems use one orchestra-

tor over N sub-agents with a maximum of 2 orchestration rounds and one re-task per round. Decentralized systems run N agents through up to 3 discussion rounds, stopping early only on true unanimity. Hybrid systems combine a generalist triage gate with the specialist panel and escalate to full discussion when the panel has no majority. Agent temperature is 0.7 within panels, 0.3 for single-doctor baselines and for orchestrator/moderator turns; reasoning models are called at effort=low with a minimum completion budget of 1,600 tokens, since reasoning tokens are drawn from the same budget and a smaller cap returns empty output.

D.3 Prompt compilation

Every agent receives the same rendered item (stem, then options as (A) ...) and the same answer schema: JSON with answer, confidence (0–100) and reason (at most 45 words). Only the system prompt differs, naming the specialty. Peer context in communicating architectures is capped at a fixed 900-character budget, balanced-truncated across peers and *independent of N* , so a panel-size effect cannot be a context-length effect. Peer ordering is deterministically rotated per (item, round) so that the first-listed speaker is decorrelated from the top-ranked specialty.

D.4 Evaluation methodology

Sample sizes. 250 items per benchmark, 180 multi-agent configurations, 63,499 episodes. **Controls.** All architectures share the item rendering, answer schema, role-prompt template and decoding parameters; only coordination logic differs. Model weights are frozen. The primary endpoint is exact match against the dataset’s own gold op-

tion key. No LLM judge appears anywhere in the evaluation path. Episodes that fail for infrastructure reasons carry an error status and are excluded from accuracy rather than scored as wrong, which would bias call-heavy architectures downward. Cost accounting is nominal: cached calls are charged at list price, so reported economics describe what a configuration would cost to run.

D.5 Information gain computation

Information gain is the Bayesian posterior variance reduction (Kim et al., 2026)

$$\Delta\mathcal{I} = \frac{1}{2} \log \frac{\text{Var}[Y \mid \mathbf{s}_{\text{pre}}]}{\text{Var}[Y \mid \mathbf{s}_{\text{post}}]}, \quad (3)$$

with $\text{Var}[Y \mid \mathbf{s}] = \hat{p}(\mathbf{s})(1 - \hat{p}(\mathbf{s}))$ and $Y \in \{0, 1\}$ the success indicator. They estimate $\hat{p}$ from $K=10$ sampled reasoning traces at temperature 0.7. In multiple-choice consultation the panel already supplies such a sample: $\hat{p}$ is the fraction of the round’s opinions that name the gold option, Laplace smoothed as $(k+1)/(n+2)$. Smoothing rather than clipping matters. A Centralized panel’s final round contains only the orchestrator, and clipping would pin $\hat{p}$ at 0.5 and force $\Delta\mathcal{I}$ negative by construction.

E Complete results grid

F Additional results

F.1 Where capability belongs inside a panel

Specialist capability matters four times more than the orchestrator’s. Replacing the whole panel with weak models costs 34.4 pp (49.2% to 14.8%). Splitting that substitution isolates where the loss comes from. Downgrading only the *orchestrator*, keeping strong specialists, costs 3.6 pp, so the panel retains 90% of the homogeneous-high span. Downgrading only the *specialists*, keeping a strong orchestrator, costs 14.8 pp, retaining 57%. A strong attending physician recovers more than half of what a weak panel loses, but a weak attending physician barely dents a strong one. The asymmetry is $4.1\times$. In a centralized clinical panel, capability should therefore be spent on the agents doing the reasoning rather than on the agent doing the auditing.

An orchestrator too weak to audit is worse than no orchestrator at all. The largest single degradation anywhere in this study is a Centralized panel with a weak model on MedQA at $P_{\text{SA}} =$

66.0%: -13.60 pp, $p = 10^{-5}$. Inspection of those episodes shows the mechanism directly. The orchestrator overrides a correct specialist majority. This is the error-amplification term in its most extreme form. It is also the counterpart to the previous result. The audit is what makes centralized coordination work, so an auditor that cannot audit converts the architecture’s strength into its worst liability.

Mixed-capability panels land between their components, with no emergent benefit. A panel of one strong, one mid and one weak specialist reaches 40.8% under Centralized and 44.0% under Decentralized, against 49.2% and 47.6% for homogeneous-high. Both sit below the strong homogeneous panel and well above the weak one, roughly where a capability-weighted average would put them. We find no medical analogue of the heterogeneous-mixing advantage reported for one model family (Kim et al., 2026), in which a low-capability orchestrator with high-capability sub-agents outperformed a homogeneous high-capability panel by 31%. In our data that configuration is the second-best Centralized cell but still 7% below homogeneous-high, directionally matching what they observe for OpenAI models, where heterogeneous centralized configurations degrade.

F.2 Where in medicine collaboration pays

Collaboration helps diagnosis most, and it is the only task type significant at both tiers. MedXpertQA labels each item as Diagnosis, Treatment, or Basic Science. Splitting the window-resident cells by label, the best panel gains +9.5 pp on Diagnosis at gpt-5-nano ($p = 0.008$) and +10.7 pp at gpt-5-mini ($p = 0.004$), against +7.1 and +8.3 pp on Treatment and +7.3 and +8.5 pp on Basic Science. Diagnosis is the only category significant at both tiers. This is the pattern clinical intuition predicts: differential diagnosis is the medical task that most resembles pooling independent perspectives, and it is where pooling pays.

The difficulty stratification is where the window shows up within a single benchmark. On MedXpertQA at gpt-5-mini, the best configuration gains +7.7 pp on easy items ($n=13$), +6.5 pp on medium ($n=62$) and +9.6 pp on hard ($n=125$). Gains are positive in every stratum for this tier, which sits inside the window; the sign flips only for tiers whose baseline falls outside it. Item diffi-




Figure 7: **Number-of-agents scaling.** Performance against panel size $n_a \in \{1, 3, 5, 7, 9\}$ for every architecture, model and benchmark. Dashed grey: single-agent baseline. Dotted: budget-matched self-consistency. Where gains occur they are captured by $n_a = 3$.

culty and configuration baseline are two views of the same quantity. The window is defined on the view that can be measured before deployment.

G Effect sizes by configuration




Figure 8: **Single doctor versus panel: performance distributions.** Box plots over all configurations of each architecture, with individual configurations overlaid. Bars are labelled SAS (single doctor), Cen, Dec, Ind and Hyb for Centralized, Decentralized, Independent and Hybrid. Annotations give the change in mean relative to the single doctor, as a percentage of that baseline. The spread within an architecture dwarfs the difference between architectures.




Figure 9: **Cost-performance trade-offs.** Mean performance against nominal cost per 1,000 questions, with SEM error bars on both axes. Each point is one configuration.




Figure 10: **Agent heterogeneity.** Performance of Centralized and Decentralized panels at $n_a=3$ on MedXpertQA under five capability placements. Strong is gpt-5-mini, weak is gpt-4.1-nano; *strong specialists* pairs a weak attending with strong specialists and *strong attending* reverses that pairing. Annotations give the change relative to the all-strong panel (dashed line). Decentralized coordination has no attending, so the two attending-capability conditions are structurally undefined there and are omitted.

Table 7: Complete results grid (1 of 4). All 180 multi-agent configurations and their single-model controls.

Tier	Benchmark	Architecture	n_a	Acc. (%)	95% CI	USD/1k
T1	MedAgentsBench	Centralized	1	20.0	[15.5, 25.4]	0.14
T1	MedAgentsBench	Centralized	3	20.4	[15.9, 25.8]	0.25
T1	MedAgentsBench	Centralized	5	20.4	[15.9, 25.8]	0.36
T1	MedAgentsBench	Centralized	7	21.6	[16.9, 27.1]	0.47
T1	MedAgentsBench	Centralized	9	19.6	[15.2, 25.0]	0.56
T1	MedAgentsBench	SAS CoT	1	20.4	[15.9, 25.8]	0.08
T1	MedAgentsBench	Decentralized	1	18.0	[13.7, 23.2]	0.05
T1	MedAgentsBench	Decentralized	3	15.6	[11.6, 20.6]	0.21
T1	MedAgentsBench	Decentralized	5	18.4	[14.1, 23.7]	0.43
T1	MedAgentsBench	Decentralized	7	17.2	[13.0, 22.4]	0.61
T1	MedAgentsBench	Decentralized	9	16.8	[12.7, 21.9]	0.88
T1	MedAgentsBench	Independent	1	18.0	[13.7, 23.2]	0.05
T1	MedAgentsBench	Independent	3	17.6	[13.4, 22.8]	0.14
T1	MedAgentsBench	Independent	5	16.4	[12.3, 21.5]	0.24
T1	MedAgentsBench	Independent	7	16.0	[12.0, 21.1]	0.33
T1	MedAgentsBench	Independent	9	16.4	[12.3, 21.5]	0.43
T1	MedAgentsBench	Self-consist.	1	20.0	[15.5, 25.4]	0.05
T1	MedAgentsBench	Self-consist.	3	18.8	[14.4, 24.1]	0.09
T1	MedAgentsBench	Self-consist.	5	19.2	[14.8, 24.5]	0.13
T1	MedAgentsBench	Self-consist.	9	18.0	[13.7, 23.2]	0.25
T1	MedAgentsBench	Self-consist.	15	18.4	[14.1, 23.7]	0.37
T1	MedAgentsBench	Hybrid	1	16.0	[12.0, 21.1]	0.05
T1	MedAgentsBench	Hybrid	3	16.0	[12.0, 21.1]	0.05
T1	MedAgentsBench	Hybrid	5	16.0	[12.0, 21.1]	0.05
T1	MedAgentsBench	Hybrid	7	16.0	[12.0, 21.1]	0.05
T1	MedAgentsBench	Hybrid	9	16.0	[12.0, 21.1]	0.05
T1	MedAgentsBench	SAS zero-shot	1	16.8	[12.7, 21.9]	0.03
T1	MedQA	Centralized	1	52.4	[46.2, 58.5]	0.14
T1	MedQA	Centralized	3	62.8	[56.7, 68.6]	0.25
T1	MedQA	Centralized	5	63.2	[57.1, 68.9]	0.36
T1	MedQA	Centralized	7	63.6	[57.5, 69.3]	0.46
T1	MedQA	Centralized	9	64.0	[57.9, 69.7]	0.57
T1	MedQA	SAS CoT	1	66.0	[59.9, 71.6]	0.08
T1	MedQA	Decentralized	1	66.4	[60.3, 72.0]	0.05
T1	MedQA	Decentralized	3	66.8	[60.7, 72.3]	0.20
T1	MedQA	Decentralized	5	67.6	[61.6, 73.1]	0.35
T1	MedQA	Decentralized	7	66.8	[60.7, 72.3]	0.56
T1	MedQA	Decentralized	9	67.2	[61.2, 72.7]	0.75
T1	MedQA	Independent	1	66.4	[60.3, 72.0]	0.05
T1	MedQA	Independent	3	66.8	[60.7, 72.3]	0.15
T1	MedQA	Independent	5	66.8	[60.7, 72.3]	0.25
T1	MedQA	Independent	7	66.4	[60.3, 72.0]	0.35
T1	MedQA	Independent	9	67.2	[61.2, 72.7]	0.45
T1	MedQA	Self-consist.	1	68.4	[62.4, 73.8]	0.05
T1	MedQA	Self-consist.	3	65.6	[59.5, 71.2]	0.10
T1	MedQA	Self-consist.	5	65.6	[59.5, 71.2]	0.14
T1	MedQA	Self-consist.	9	66.0	[59.9, 71.6]	0.26
T1	MedQA	Self-consist.	15	65.6	[59.5, 71.2]	0.40
T1	MedQA	Hybrid	1	65.6	[59.5, 71.2]	0.05
T1	MedQA	Hybrid	3	65.6	[59.5, 71.2]	0.05
T1	MedQA	Hybrid	5	65.6	[59.5, 71.2]	0.05
T1	MedQA	Hybrid	7	65.6	[59.5, 71.2]	0.05
T1	MedQA	Hybrid	9	65.6	[59.5, 71.2]	0.05
T1	MedQA	SAS zero-shot	1	66.4	[60.3, 72.0]	0.03
T1	MedXpertQA	Centralized	1	19.2	[14.8, 24.5]	0.17
T1	MedXpertQA	Centralized	3	16.0	[12.0, 21.1]	0.28
T1	MedXpertQA	Centralized	5	16.0	[12.0, 21.1]	0.40
T1	MedXpertQA	Centralized	7	17.6	[13.4, 22.8]	0.52
T1	MedXpertQA	Centralized	9	14.8	[10.9, 19.7]	0.62
T1	MedXpertQA	SAS CoT	1	15.6	[11.6, 20.6]	0.09
T1	MedXpertQA	Decentralized	1	16.0	[12.0, 21.1]	0.05
T1	MedXpertQA	Decentralized	3	15.2	[11.3, 20.2]	0.25

Table 8: Complete results grid (2 of 4). All 180 multi-agent configurations and their single-model controls.

Tier	Benchmark	Architecture	n_a	Acc. (%)	95% CI	USD/1k
T1	MedXpertQA	Decentralized	5	16.8	[12.7, 21.9]	0.50
T1	MedXpertQA	Decentralized	7	16.0	[12.0, 21.1]	0.74
T1	MedXpertQA	Decentralized	9	16.4	[12.3, 21.5]	1.05
T1	MedXpertQA	Independent	1	15.8	[12.9, 19.3]	0.05
T1	MedXpertQA	Independent	3	15.8	[12.9, 19.3]	0.16
T1	MedXpertQA	Independent	5	16.0	[12.0, 21.1]	0.27
T1	MedXpertQA	Independent	7	16.0	[12.0, 21.1]	0.38
T1	MedXpertQA	Independent	9	16.0	[12.0, 21.1]	0.48
T1	MedXpertQA	Self-consist.	1	16.0	[12.0, 21.1]	0.06
T1	MedXpertQA	Self-consist.	3	16.4	[12.3, 21.5]	0.10
T1	MedXpertQA	Self-consist.	5	16.0	[12.0, 21.1]	0.14
T1	MedXpertQA	Self-consist.	9	15.6	[11.6, 20.6]	0.27
T1	MedXpertQA	Self-consist.	15	16.0	[12.0, 21.1]	0.40
T1	MedXpertQA	Hybrid	1	18.0	[13.7, 23.2]	0.05
T1	MedXpertQA	Hybrid	3	18.0	[13.7, 23.2]	0.05
T1	MedXpertQA	Hybrid	5	18.0	[13.7, 23.2]	0.05
T1	MedXpertQA	Hybrid	7	18.0	[13.7, 23.2]	0.06
T1	MedXpertQA	Hybrid	9	18.0	[13.7, 23.2]	0.06
T1	MedXpertQA	SAS zero-shot	1	16.4	[12.3, 21.5]	0.03
T2	MedAgentsBench	Centralized	1	32.8	[27.3, 38.8]	0.62
T2	MedAgentsBench	Centralized	3	34.4	[28.8, 40.5]	0.93
T2	MedAgentsBench	Centralized	5	32.8	[27.3, 38.8]	1.25
T2	MedAgentsBench	Centralized	7	32.4	[26.9, 38.4]	1.59
T2	MedAgentsBench	Centralized	9	32.8	[27.3, 38.8]	1.91
T2	MedAgentsBench	SAS CoT	1	32.4	[26.9, 38.4]	0.19
T2	MedAgentsBench	Decentralized	1	33.6	[28.0, 39.7]	0.16
T2	MedAgentsBench	Decentralized	3	34.0	[28.4, 40.1]	0.73
T2	MedAgentsBench	Decentralized	5	33.2	[27.7, 39.3]	1.34
T2	MedAgentsBench	Decentralized	7	34.0	[28.4, 40.1]	2.11
T2	MedAgentsBench	Decentralized	9	34.4	[28.8, 40.5]	2.92
T2	MedAgentsBench	Independent	1	33.6	[28.0, 39.7]	0.16
T2	MedAgentsBench	Independent	3	34.8	[29.2, 40.9]	0.48
T2	MedAgentsBench	Independent	5	34.8	[29.2, 40.9]	0.79
T2	MedAgentsBench	Independent	7	33.2	[27.7, 39.3]	1.12
T2	MedAgentsBench	Independent	9	32.4	[26.9, 38.4]	1.44
T2	MedAgentsBench	Self-consist.	1	33.2	[27.7, 39.3]	0.18
T2	MedAgentsBench	Self-consist.	3	32.0	[26.5, 38.0]	0.50
T2	MedAgentsBench	Self-consist.	5	33.6	[28.0, 39.7]	0.83
T2	MedAgentsBench	Self-consist.	9	32.8	[27.3, 38.8]	1.49
T2	MedAgentsBench	Self-consist.	15	31.2	[25.8, 37.2]	2.47
T2	MedAgentsBench	Hybrid	1	32.4	[26.9, 38.4]	0.26
T2	MedAgentsBench	Hybrid	3	33.2	[27.7, 39.3]	0.49
T2	MedAgentsBench	Hybrid	5	34.0	[28.4, 40.1]	0.73
T2	MedAgentsBench	Hybrid	7	32.8	[27.3, 38.8]	0.94
T2	MedAgentsBench	Hybrid	9	32.8	[27.3, 38.8]	1.16
T2	MedAgentsBench	SAS zero-shot	1	33.2	[27.7, 39.3]	0.13
T2	MedQA	Centralized	1	86.4	[81.6, 90.1]	0.47
T2	MedQA	Centralized	3	88.8	[84.3, 92.1]	0.67
T2	MedQA	Centralized	5	89.6	[85.2, 92.8]	0.95
T2	MedQA	Centralized	7	88.4	[83.8, 91.8]	1.23
T2	MedQA	Centralized	9	86.0	[81.2, 89.8]	1.51
T2	MedQA	SAS CoT	1	88.0	[83.4, 91.5]	0.16
T2	MedQA	Decentralized	1	87.6	[82.9, 91.1]	0.13
T2	MedQA	Decentralized	3	88.8	[84.3, 92.1]	0.50
T2	MedQA	Decentralized	5	89.2	[84.7, 92.5]	0.86
T2	MedQA	Decentralized	7	88.8	[84.3, 92.1]	1.26
T2	MedQA	Decentralized	9	89.6	[85.2, 92.8]	1.75
T2	MedQA	Independent	1	87.6	[82.9, 91.1]	0.13
T2	MedQA	Independent	3	88.4	[83.8, 91.8]	0.39
T2	MedQA	Independent	5	89.6	[85.2, 92.8]	0.65
T2	MedQA	Independent	7	87.6	[82.9, 91.1]	0.90
T2	MedQA	Independent	9	88.8	[84.3, 92.1]	1.16

Table 9: Complete results grid (3 of 4). All 180 multi-agent configurations and their single-model controls.

Tier	Benchmark	Architecture	n_a	Acc. (%)	95% CI	USD/1k
T2	MedQA	Self-consist.	1	84.4	[79.4, 88.4]	0.15
T2	MedQA	Self-consist.	3	86.8	[82.0, 90.4]	0.41
T2	MedQA	Self-consist.	5	88.0	[83.4, 91.5]	0.67
T2	MedQA	Self-consist.	9	87.6	[82.9, 91.1]	1.21
T2	MedQA	Self-consist.	15	88.4	[83.8, 91.8]	2.01
T2	MedQA	Hybrid	1	87.6	[82.9, 91.1]	0.17
T2	MedQA	Hybrid	3	87.2	[82.5, 90.8]	0.26
T2	MedQA	Hybrid	5	88.8	[84.3, 92.1]	0.35
T2	MedQA	Hybrid	7	87.2	[82.5, 90.8]	0.43
T2	MedQA	Hybrid	9	87.6	[82.9, 91.1]	0.53
T2	MedQA	SAS zero-shot	1	87.2	[82.5, 90.8]	0.09
T2	MedXpertQA	Centralized	1	31.4	[27.5, 35.6]	0.73
T2	MedXpertQA	Centralized	3	36.0	[31.9, 40.3]	1.02
T2	MedXpertQA	Centralized	5	35.0	[30.9, 39.3]	1.38
T2	MedXpertQA	Centralized	7	36.4	[30.7, 42.5]	1.80
T2	MedXpertQA	Centralized	9	37.2	[31.4, 43.3]	2.17
T2	MedXpertQA	SAS CoT	1	32.0	[26.5, 38.0]	0.20
T2	MedXpertQA	Decentralized	1	34.0	[28.4, 40.1]	0.19
T2	MedXpertQA	Decentralized	3	38.0	[32.2, 44.2]	0.91
T2	MedXpertQA	Decentralized	5	38.0	[32.2, 44.2]	1.66
T2	MedXpertQA	Decentralized	7	38.0	[32.2, 44.2]	2.46
T2	MedXpertQA	Decentralized	9	37.2	[31.4, 43.3]	3.28
T2	MedXpertQA	Independent	1	31.8	[27.9, 36.0]	0.19
T2	MedXpertQA	Independent	3	32.6	[28.6, 36.8]	0.56
T2	MedXpertQA	Independent	5	33.8	[29.8, 38.1]	0.93
T2	MedXpertQA	Independent	7	32.8	[28.8, 37.0]	1.30
T2	MedXpertQA	Independent	9	32.6	[28.6, 36.8]	1.67
T2	MedXpertQA	Self-consist.	1	32.8	[27.3, 38.8]	0.20
T2	MedXpertQA	Self-consist.	3	36.0	[30.3, 42.1]	0.57
T2	MedXpertQA	Self-consist.	5	34.8	[29.2, 40.9]	0.94
T2	MedXpertQA	Self-consist.	9	35.2	[29.5, 41.3]	1.70
T2	MedXpertQA	Self-consist.	15	34.4	[28.8, 40.5]	2.81
T2	MedXpertQA	Hybrid	1	33.3	[27.8, 39.4]	0.29
T2	MedXpertQA	Hybrid	3	34.4	[28.8, 40.5]	0.60
T2	MedXpertQA	Hybrid	5	35.2	[29.5, 41.3]	0.89
T2	MedXpertQA	Hybrid	7	33.6	[28.0, 39.7]	1.18
T2	MedXpertQA	Hybrid	9	33.2	[27.7, 39.3]	1.56
T2	MedXpertQA	SAS zero-shot	1	32.4	[26.9, 38.4]	0.15
T3	MedAgentsBench	Centralized	1	44.8	[38.8, 51.0]	1.54
T3	MedAgentsBench	Centralized	3	44.4	[38.4, 50.6]	2.40
T3	MedAgentsBench	Centralized	5	44.8	[38.8, 51.0]	3.51
T3	MedAgentsBench	Centralized	7	46.0	[39.9, 52.2]	4.60
T3	MedAgentsBench	Centralized	9	45.6	[39.5, 51.8]	5.76
T3	MedAgentsBench	SAS CoT	1	44.4	[38.4, 50.6]	0.68
T3	MedAgentsBench	Decentralized	1	44.4	[38.4, 50.6]	0.56
T3	MedAgentsBench	Decentralized	3	44.8	[38.8, 51.0]	2.11
T3	MedAgentsBench	Decentralized	5	45.2	[39.1, 51.4]	3.88
T3	MedAgentsBench	Decentralized	7	46.4	[40.3, 52.6]	5.99
T3	MedAgentsBench	Decentralized	9	46.8	[40.7, 53.0]	7.80
T3	MedAgentsBench	Independent	1	44.4	[38.4, 50.6]	0.55
T3	MedAgentsBench	Independent	3	44.4	[38.4, 50.6]	1.61
T3	MedAgentsBench	Independent	5	44.0	[38.0, 50.2]	2.67
T3	MedAgentsBench	Independent	7	43.6	[37.6, 49.8]	3.72
T3	MedAgentsBench	Independent	9	44.8	[38.8, 51.0]	4.77
T3	MedAgentsBench	Self-consist.	1	43.6	[37.6, 49.8]	0.62
T3	MedAgentsBench	Self-consist.	3	44.4	[38.4, 50.6]	1.73
T3	MedAgentsBench	Self-consist.	5	44.4	[38.4, 50.6]	2.84
T3	MedAgentsBench	Self-consist.	9	44.0	[38.0, 50.2]	5.12
T3	MedAgentsBench	Self-consist.	15	44.4	[38.4, 50.6]	8.43
T3	MedAgentsBench	Hybrid	1	46.0	[39.9, 52.2]	0.62
T3	MedAgentsBench	Hybrid	3	45.2	[39.1, 51.4]	0.78
T3	MedAgentsBench	Hybrid	5	44.8	[38.8, 51.0]	0.99

Table 10: Complete results grid (4 of 4). All 180 multi-agent configurations and their single-model controls.

Tier	Benchmark	Architecture	n_a	Acc. (%)	95% CI	USD/1k
T3	MedAgentsBench	Hybrid	7	45.2	[39.1, 51.4]	1.09
T3	MedAgentsBench	Hybrid	9	45.6	[39.5, 51.8]	1.25
T3	MedAgentsBench	SAS zero-shot	1	46.0	[39.9, 52.2]	0.39
T3	MedQA	Centralized	1	93.2	[89.4, 95.7]	1.07
T3	MedQA	Centralized	3	91.6	[87.5, 94.4]	2.00
T3	MedQA	Centralized	5	92.4	[88.4, 95.1]	2.84
T3	MedQA	Centralized	7	92.4	[88.4, 95.1]	3.80
T3	MedQA	Centralized	9	92.4	[88.4, 95.1]	4.66
T3	MedQA	SAS CoT	1	91.6	[87.5, 94.4]	0.57
T3	MedQA	Decentralized	1	92.8	[88.9, 95.4]	0.44
T3	MedQA	Decentralized	3	92.4	[88.4, 95.1]	1.50
T3	MedQA	Decentralized	5	92.8	[88.9, 95.4]	2.62
T3	MedQA	Decentralized	7	92.4	[88.4, 95.1]	3.70
T3	MedQA	Decentralized	9	92.4	[88.4, 95.1]	4.70
T3	MedQA	Independent	1	92.4	[88.4, 95.1]	0.43
T3	MedQA	Independent	3	91.6	[87.5, 94.4]	1.29
T3	MedQA	Independent	5	91.6	[87.5, 94.4]	2.15
T3	MedQA	Independent	7	91.6	[87.5, 94.4]	3.01
T3	MedQA	Independent	9	92.0	[88.0, 94.8]	3.86
T3	MedQA	Self-consist.	1	92.0	[88.0, 94.8]	0.52
T3	MedQA	Self-consist.	3	92.8	[88.9, 95.4]	1.41
T3	MedQA	Self-consist.	5	93.2	[89.4, 95.7]	2.30
T3	MedQA	Self-consist.	9	94.0	[90.3, 96.3]	4.15
T3	MedQA	Self-consist.	15	93.6	[89.9, 96.0]	6.78
T3	MedQA	Hybrid	1	93.6	[89.9, 96.0]	0.43
T3	MedQA	Hybrid	3	93.2	[89.4, 95.7]	0.48
T3	MedQA	Hybrid	5	93.2	[89.4, 95.7]	0.52
T3	MedQA	Hybrid	7	93.2	[89.4, 95.7]	0.56
T3	MedQA	Hybrid	9	93.2	[89.4, 95.7]	0.60
T3	MedQA	SAS zero-shot	1	94.4	[90.8, 96.6]	0.27
T3	MedXpertQA	Centralized	1	46.8	[40.7, 53.0]	1.53
T3	MedXpertQA	Centralized	3	49.2	[43.1, 55.4]	2.48
T3	MedXpertQA	Centralized	5	48.4	[42.3, 54.6]	3.68
T3	MedXpertQA	Centralized	7	46.0	[39.9, 52.2]	4.89
T3	MedXpertQA	Centralized	9	46.4	[40.3, 52.6]	6.07
T3	MedXpertQA	SAS CoT	1	41.6	[35.7, 47.8]	0.72
T3	MedXpertQA	Decentralized	1	45.6	[39.5, 51.8]	0.58
T3	MedXpertQA	Decentralized	3	47.6	[41.5, 53.8]	2.33
T3	MedXpertQA	Decentralized	5	48.0	[41.9, 54.2]	4.39
T3	MedXpertQA	Decentralized	7	46.8	[40.7, 53.0]	6.65
T3	MedXpertQA	Decentralized	9	46.0	[39.9, 52.2]	9.14
T3	MedXpertQA	Independent	1	41.6	[37.4, 46.0]	0.59
T3	MedXpertQA	Independent	3	46.4	[40.3, 52.6]	1.71
T3	MedXpertQA	Independent	5	47.2	[41.1, 53.4]	2.84
T3	MedXpertQA	Independent	7	46.0	[39.9, 52.2]	3.97
T3	MedXpertQA	Independent	9	46.4	[40.3, 52.6]	5.11
T3	MedXpertQA	Self-consist.	1	42.4	[36.4, 48.6]	0.67
T3	MedXpertQA	Self-consist.	3	43.2	[37.2, 49.4]	1.83
T3	MedXpertQA	Self-consist.	5	43.2	[37.2, 49.4]	3.00
T3	MedXpertQA	Self-consist.	9	46.0	[39.9, 52.2]	5.41
T3	MedXpertQA	Self-consist.	15	47.6	[41.5, 53.8]	8.87
T3	MedXpertQA	Hybrid	1	45.2	[39.1, 51.4]	0.61
T3	MedXpertQA	Hybrid	3	44.4	[38.4, 50.6]	0.71
T3	MedXpertQA	Hybrid	5	44.8	[38.8, 51.0]	0.87
T3	MedXpertQA	Hybrid	7	44.4	[38.4, 50.6]	0.93
T3	MedXpertQA	Hybrid	9	44.4	[38.4, 50.6]	1.20
T3	MedXpertQA	SAS zero-shot	1	43.2	[37.2, 49.4]	0.46

Table 11: Capability ladder, measured on 40 items per benchmark with single-doctor CoT.

Tier	Model	Effort	MedXpertQA	MedQA	USD / 1k items
T1	gpt-4.1-nano	—	25.0%	75.0%	0.09
T2	gpt-5-nano	low	37.5%	95.0%	0.17
T3	gpt-5-mini	low	55.0%	97.5%	0.64

Table 12: Agent heterogeneity at $n_a=3$ on MedXpertQA. High capability is gpt-5-mini, low is gpt-4.1-nano. “Span retained” is the fraction of the homogeneous-high minus homogeneous-low range that the configuration preserves. Decentralized coordination has no orchestrator, so those conditions are structurally undefined.

Capability placement	Centralized		Decentralized	
	accuracy	span retained	accuracy	span retained
Homogeneous high	49.2%	100%	47.6%	100%
Low orchestrator, high agents	45.6%	90%	—	—
High orchestrator, low agents	34.4%	57%	—	—
Mixed-capability panel	40.8%	76%	44.0%	88%
Homogeneous low	14.8%	0%	17.2%	0%

Table 13: Collaboration gain by clinical task type on MedXpertQA, for the two tiers inside the window. Gain is the best panel configuration over the single doctor on the same items; p from paired McNemar.

Tier	Task type	n	single doctor	best panel	gain
gpt-5-nano	Diagnosis	84	33.3%	42.9%	+9.52 pp**
	Treatment	84	35.7%	42.9%	+7.14 pp
	Basic Science	82	26.8%	34.1%	+7.32 pp
gpt-5-mini	Diagnosis	84	38.1%	48.8%	+10.71 pp**
	Treatment	84	44.0%	52.4%	+8.33 pp*
	Basic Science	82	42.7%	51.2%	+8.54 pp*

* $p < 0.05$, ** $p < 0.01$




Figure 11: **Effect size for every configuration, ranked.** Collaboration gain with 95% confidence intervals; each row is one architecture at its best panel size within a difficulty stratum, and rows run left to right in descending order of effect. Amber rows have a single-doctor baseline inside the 25–50% window, grey rows outside it. Filled markers are significant at $p < 0.05$, hollow are not. The ranking is what carries the claim: amber occupies the positive tail and grey the negative one, and every significant row on the positive side is amber while the only significant row on the negative side is grey.